

Chapter 10

Analytic Functions & CTE's

DRAFT

Contents

1	Analytic Functions	167
2	Using Analytic Functions with Transaction Data	174
3	Common Table Expressions (“CTE”)	176
4	CTEs with the transaction data	178

DRAFT

1 Analytic Functions

- Analytic (sometimes called “window” or “partition” functions) functions were designed to simplify many common complex joins.
- There are a few different use cases that they can greatly simplify.
- As an example, lets consider the case of computing the *percentage* of traffic which pays by cash by hour and plaza in the inbound direction on November 10th, 2016. To do this we would need to take our original data and then join it against the correct sum. In other words we want to return 24 columns for each plaza and the sum (vertically, across hour) should be equal to 1.
- In order to do this calculation we need to join our original data back onto the proper sum, as can be seen in the query below.

```
select
  plaza
  , hr
  , vehiclescash::float / ALLVech as pctperhr
  , vehiclescash
  , ALLVech
from
  (select plaza, hr, vehiclescash
   from cls.mta
   where mtadt = '2016-11-10'
   and direction = 'I') as lhs
left join
  (select plaza, sum(vehiclescash) as ALLVech
   from cls.mta
   where mtadt = '2016-11-10'
   and direction = 'I'
   group by 1) as rhs
using(plaza)
order by 1,2;
```

plaza	hr	pctperhr	vehiclescash	allvech
-----	----	-----	-----	-----
1	0	0.0230918	167	7232
1	1	0.017146	124	7232
1	2	0.00954093	69	7232
1	3	0.00940265	68	7232
1	4	0.0199115	144	7232
[...]				

but this construct is a bit cumbersome.

- For another example, consider wanting to create a rolling sum over each plaza day for the number of cars which use cash in the inbound direction.

```

select
    lhs.plaza, lhs.mtadt, lhs.hr, sum( rhs.vehiclesscash ) as cum_sum
from
    (select plaza, hr, mtadt
     from cls.mta where direction = 'I') as lhs
left join
    (select plaza, hr, mtadt, vehiclesscash
     from cls.mta where direction = 'I') as rhs
on
    lhs.plaza = rhs.plaza
    and lhs.hr >= rhs.hr
    and lhs.mtadt = rhs.mtadt
group by lhs.plaza, lhs.mtadt, lhs.hr
order by 1,2,3

```

plaza	mtadt	hr	cum_sum
1	2010-01-01	0	474
1	2010-01-01	1	1191
1	2010-01-01	2	1855
1	2010-01-01	3	2450
1	2010-01-01	4	2997
[...]			

- These two examples have a set of common properties: we need to aggregate over our table while returning the original table. Doing this using the techniques we've seen in the past is cumbersome, so we can use Analytic (sometimes called window or partition functions) to solve them.
- Analytic functions use the following syntax:

```

function () over(
    partition by _____
    order by _____
    <WINDOW FRAME CLAUSE>
)

```

function can be one of any of our standard aggregate functions (SUM, COUNT, MAX, MIN, AVG) as well as a number of functions that can only be used as analytic functions.

There are a few pieces of the syntax:

1. The OVER() clause: This tells the database to expect a window function, rather than a standard aggregate function. This is required when using analytic functions.
2. The PARTITION BY clause: This clause tells the database how to break up the data. In other words, it is similar to a GROUP BY in that it tells the database that rows with the same values should be treated as a single entity or partition. The PARTITION BY clause is optional.
3. The ORDER BY clause: This clause works just as an ORDER BY in a normal SQL query works. It tells the database how to sort the data within each partition. The ORDER BY clause is optional. If an ORDER BY clause is present then the function is calculated in a running

fashion – e.g. as a running sum from the start of the partition to the current row.

4. The WINDOW FRAME clause defines the region over which the function is calculated. It takes on a number of different forms though the most common is the rows between syntax:

ROWS BETWEEN _____ AND _____

the blanks would take on some of the following values:

- UNBOUNDED PRECEDING: from the start of the partition
- UNBOUNDED FOLLOWING: to the end of the partition
- XX PRECEDING: XX rows preceding (inclusive)
- XX FOLLOWING: XX rows following (inclusive)
- CURRENT ROW: the current row

In other words we can use this syntax to easily compute things like hourly moving averages or just smoothing.

- Lets use analytics functions to solve the two problems at the start of this section. To solve the first one we can do the following:

```
select
  plaza
  , hr
  , vehiclesscash::float/sum(vehiclesscash) over(partition by plaza) as pctperhr
  , vehiclesscash
  , sum( vehiclesscash) over(partition by plaza) as ALLVech
from
  cls.mta
where
  mtdt = '2016-11-10' and direction = 'I'
order by 1,2;
```

plaza	hr	pctperhr	vehiclesscash	allvech
-----	----	-----	-----	-----
1	0	0.0230918	167	7232
1	1	0.017146	124	7232
1	2	0.00954093	69	7232
1	3	0.00940265	68	7232
1	4	0.0199115	144	7232
[...]				

the OVER clause, which modifies the SUM function, tells the database that it is going to be computing an aggregate function, but *without the aggregation*. In other words, it will return the same value for each row.

- To be clear on what this is doing, let's consider only a single plaza (#1) and look at the hourly data for that day, including the analytic function:

```

select
  plaza
  , hr
  , vehiclesscash
  , sum(vehiclesscash) over(partition by plaza) as totalcars
from
  cls.mta
where
  mtadt = '2016-11-10'
  and direction = 'I'
  and plaza = 1;

```

plaza	hr	vehiclesscash	totalcars
1	0	167	7232
1	1	124	7232
1	2	69	7232
1	3	68	7232
1	4	144	7232
1	5	215	7232
1	6	281	7232
1	7	336	7232
1	8	329	7232
1	9	304	7232
1	10	344	7232
1	11	286	7232
1	12	308	7232
1	13	375	7232
1	14	361	7232
1	15	471	7232
1	16	450	7232
1	17	451	7232
1	18	420	7232
1	19	446	7232
1	20	366	7232
1	21	322	7232
1	22	296	7232
1	23	299	7232

The sum of vehiclesscash on this subset is 7,232 – the exact number returned by the analytic function in the total cars column.

- To solve the cumulative sum problem we can use an analytic function in the following manner:

```

select
    plaza, mtadt, hr,
    sum(vehiclesscash) over(
        partition by plaza, mtadt
        order by hr
        rows between unbounded preceding and current row) as cum_sum
from
    cls.mta
where
    direction = 'I'

```

plaza	mtadt	hr	cum_sum
1	2010-01-01	0	474
1	2010-01-01	1	1191
1	2010-01-01	2	1855
1	2010-01-01	3	2450
1	2010-01-01	4	2997
[...]			

- To get more insight into the specifics of how analytic functions work, consider the following table:

GRP	ORD	NM	C0	C1	C2	C3
1	1	5	65	33	5	11
1	2	6	65	33	11	21
1	3	10	65	33	21	28
1	4	12	65	33	33	22
2	1	12	65	32	12	22
2	2	10	65	32	22	28
2	3	6	65	32	28	20
2	4	4	65	32	32	10

In this table the raw data is GRP, ORD and NM. In order to create columns C1, C2 and C3 we use the following syntax:

```

SUM(NM) OVER( ) as C0
SUM(NM) OVER( PARTITION BY GRP ) as C1
SUM(NM) OVER( PARTITION BY GRP ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as C1
SUM(NM) OVER( PARTITION BY GRP ORDER BY ORD ASC) as C2
SUM(NM) OVER( PARTITION BY GRP ORDER BY ORD ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING) AS C3

```

- **NOTE:** The default behavior of different analytic functions when using different sets of arguments can lead to issues. I have a few cases memorized, but I'd recommend being as inclusive as possible with the arguments. In the example above, C1 is presented twice, both would return the same numbers, but the second makes it more clear as to what is happening.
- All of the aggregate functions we used in the past (MAX, MIN, COUNT and AVG) can be used with analytic functions. For example, to create a moving average based on the last 4 hours of data in the MTA dataset we could do the the following. Note that the ROWS BETWEEN function is inclusive, so it will average over four hours in the below.

```

select
    plaza
    , direction
    , hr
    , mtadt
    , vehiclescash + vehiclesez as totalcars
    , avg(vehiclescash + vehiclesez)
      OVER(PARTITION by plaza, direction
           ORDER BY mtadt asc, hr asc
           ROWS BETWEEN 3 PRECEDING AND CURRENT ROW) as hravg
from
    cls.mta;

```

plaza	direction	hr	mtadt	totalcars	hravg
-----	-----	---	-----	-----	-----
1	I	0	2010-01-01	889	889
1	I	1	2010-01-01	1419	1154
1	I	2	2010-01-01	1223	1177
1	I	3	2010-01-01	1075	1151.5
1	I	4	2010-01-01	945	1165.5
[...]					

- If we use a WINDOW FRAME clause without an ORDER BY then the row order returned is arbitrary.
- The only aggregate function not allowed is COUNT(DISTINCT) which cannot be used with the OVER() clause.
- There are a number of non-aggregate functions that can be used as analytic functions:
 1. LAG() and LEAD(): These functions return the value of a column from a preceding or following row.
 2. FIRST_VALUE(), LAST_VALUE() and NTH_VALUE(): These function return the first, last, or more generally, the nth value within a partition. Note that the NTH_VALUE function takes not only a column name, but a positional argument starting at 1. It will return null if there aren't enough values within the partition.
 3. NTILE(): This function handles percentiles.
 4. ROW_NUMBER(): This function returns the row number based on the criteria established in the clause. Note that the function ROW_NUMBER() fails without an OVER clause.
 5. RANK(): Returns the rank of a particular observation
 6. DENSE_RANK(): Returns the dense rank of a particular observation.
- The commands ROW_NUMBER, RANK and DENSE_RANK behave similarly when the data being sorted is unique. If the data is not unique, however, these commands behave differently, as demonstrated in the Table 10.1
- These functions allow us to easily answer a number of questions (without using JOIN). Such as what is the correlation between the absolute (nominal) change in vehicles paying cash and the absolute (nominal) change in vehicles paying by EZ-pass?

ID	ROW_NUMBER	RANK	DENSE_RANK
1	1	1	1
1	2	1	1
1	3	1	1
1	4	1	1
2	5	5	2
2	6	5	2
3	7	7	3

Table 10.1: Ranking function differences when ordering by the ID columns

```

SELECT
    CORR( CashDiff, EZDiff ) as DiffCor
FROM
    (SELECT
        LAG(vehiclesEZ) OVER(PARTITION BY plaza
            ORDER BY mtadt ASC, hr asc)
        - vehiclesEZ AS EZDiff
    , LAG(vehiclesscash) OVER(PARTITION BY plaza
        ORDER BY mtadt ASC, hr asc)
        - vehiclesscash AS CashDiff
    FROM
        cls.mta ) as innerQ;

diffcor
-----
0.744516

```

- These types of functions are evaluated *after* with SELECT after GROUP, JOIN, WHERE and HAVING. This means that you can't refer to them within those functions. If you want to filter on a window function it must be contained within a subquery.
- An important caveat when using these functions, as said from the documentation (emphasis mine):

By default, if ORDER BY is supplied then the frame consists of all rows from the start of the partition up through the current row, **plus any following rows that are equal to the current row** according to the ORDER BY clause. When ORDER BY is omitted the default frame consists of all rows in the partition.

This is weird:

```

select
  plaza, mtadt, hr, direction
  , vehiclesez
  , sum(vehiclesez)
    over(partition by plaza order by mtadt, hr
         rows between unbounded preceding
                and current row) as runningS1
  , sum(vehiclesez)
    over(partition by plaza order by mtadt, hr ) as runningS2
from
  cls.mta
where plaza = 1 and mtadt = '2010-01-01' and hr < 3;

```

plaza	mtadt	hr	direction	vehiclesez	runningS1	runningS2
1	2010-01-01	0	I	415	415	801
1	2010-01-01	0	O	386	801	801
1	2010-01-01	1	I	702	1503	2037
1	2010-01-01	1	O	534	2037	2037
1	2010-01-01	2	I	559	2596	3106
[...]						

This is weird because when you omit ROWS BETWEEN, the running sum is computed as if rows which have similar values in the partition are the same. The same query however, with a ROWS BETWEEN clause computes a running sum while ignoring the duplicate rows.

2 Using Analytic Functions with Transaction Data

- In this section we return to trying to understand the revenue behavior of our soap transaction data. Just like before we are going to use the notion of a cohort to help our analysis.
- Consider the data in Table 9.1 which contains information on users who were making certain transactions.
- Let's begin by calculating the revenue per user by locale and also by their install time period.
- We didn't even attempt this in the previous section because it involved so many joins! Using Analytic functions allows us to skip many of those issues!

```

select
  cohort
  , locale
  , count(distinct userid) as numusers
  , sum(case when trans_dt::date <= (cohort + '1 month'::interval)::date
        then amt else 0 end ) as mon_0_amt
  , sum(case when trans_dt::date <= (cohort + '2 month'::interval)::date
        then amt else 0 end ) as mon_1_amt
  , sum(case when trans_dt::date <= (cohort + '3 month'::interval)::date
        then amt else 0 end ) as mon_2_amt
from
  ( select
    first_value(locale) over(partition by userid order by trans_dt asc ) as locale
    , date_trunc('month', first_value(trans_dt)
      over(partition by userid order by trans_dt asc))::date as cohort
    , amt, userid, trans_dt
  from
    cls.trans ) as innerQ
GROUP BY 1,2

```

cohort	locale	numusers	mon_0_amt	mon_1_amt	mon_2_amt
2016-01-01	Canada	4706	162262	167287	223375
2016-01-01	Mexico	3920	186854	190914	224451
2016-01-01	U.S.	12676	542198	599027	637072
2016-02-01	Canada	4334	147535	152976	206147
2016-02-01	Mexico	3602	171472	175171	208097
[...]					

- Let's calculate the percentage of revenue that each transaction represents for each userid (how would we do this without Analytic Functions?):

```

select
  userid, trans_dt, amt
  , amt/sum( amt) over(partition by userid) as pct
from
  cls.trans

```

userid	trans_dt	amt	pct
1	2016-05-09	23.98	1
2	2018-08-25	12.99	1
3	2017-03-05	43.16	0.5
3	2017-04-05	43.16	0.5
4	2016-02-28	59.95	1
[...]			

- I want to calculate the percentage likelihood that a person who has made X purchases makes another one. There are a number of different ways that this can be done, but we can use analytic functions:

```

select
    transNum
    , sum( case when transNum = totalTrans then 1 else 0 end)::float
    / count(1) as pct
    , count(1) as numerator
from
    (select
        row_number() over(partition by userid order by trans_dt) as transNum
        , count(1) over(partition by userid) as totalTrans
    from
        cls.trans) as innerQ
group by 1
order by 1;

```

transnum	pct	numerator
1	0.529178	574289
2	0.538423	270388
3	0.58621	124805
4	0.644172	51643
5	0.686276	18376

[...]

3 Common Table Expressions (“CTE”)

- A relatively new piece of SQL syntax is WITH, which allows for tables to be defined and used repeatedly within a query. These are called Common Table Expressions. CTE are incredibly powerful ways of writing queries, but they can come with significant downsides (as we will discuss later).

From PostgreSQL’s documentation:

A useful property of WITH queries is that they are evaluated only once per execution of the parent query, even if they are referred to more than once by the parent query or sibling WITH queries. Thus, expensive calculations that are needed in multiple places can be placed within a WITH query to avoid redundant work. Another possible application is to prevent unwanted multiple evaluations of functions with side-effects. However, the other side of this coin is that the optimizer is less able to push restrictions from the parent query down into a WITH query than an ordinary sub-query. The WITH query will generally be evaluated as written, without suppression of rows that the parent query might discard afterwards. (But, as mentioned above, evaluation might stop early if the reference(s) to the query demand only a limited number of rows.)

- The motivation for CTE is that they can increase readability in query by defining a table at the start of your query which only exists for the duration of the query.
- The WITH clause is used to start a CTE and it basically sets up a derived table that can be used in the query.
- Consider the following example:

```
with only_inbound as (select * from cls.mta where direction = 'I')

select * from only_inbound limit 100;
```

plaza	mtadt	hr	direction	vehiclesez	vehiclesscash
2	2013-10-14	16	I	2469	336
2	2013-10-14	17	I	2853	425
2	2013-10-14	18	I	2575	394
2	2013-10-14	19	I	2422	344
2	2013-10-14	20	I	1989	339
[...]					

The basic syntax of the query is that we define a table via a query at the start using a WITH clause. This table does not have a schema and can only be referenced within that query.

- We can use CTEs with multiple queries by separating them with commas:

```
with
  only_inbound as (select * from cls.mta where direction = 'I')
  , only_outbound as (select * from cls.mta where direction = 'O')

select
  plaza, mtadt, hr
  , only_inbound.vehiclesez as inbound_ez
  , only_outbound.vehiclesez as outbound_ez
from
  only_inbound
join
  only_outbound
using( plaza, mtadt, hr)

limit 10;
```

plaza	mtadt	hr	inbound_ez	outbound_ez
2	2013-10-14	17	2853	2116
2	2013-10-13	11	2960	2081
2	2013-10-12	17	2847	2433
2	2013-10-10	1	189	177
2	2013-10-10	3	118	140
[...]				

- Where I find CTEs to be useful is when there are multiple layers of logic that need to be implemented. By using a CTE I can break up that application logic into separate pieces that are easier to read.

4 CTEs with the transaction data

- To use the WITH clause you specify a table name and then use AS. It is done before the SELECT in the query. For example, the following creates a table that only looks at unit transactions from the United States. We then use this figure out the average order value of these transactions:

```
with USUnits as (  
    select * from cls.trans  
    where locale = 'U.S.' and type = 'Units')  
  
select avg(amt) as AOV from USUnits;  
  
      aov  
-----  
43.3045
```

- Consider the following, more useful, example which creates an LTV dataset which has the first value of the local of purchase.

```
with LTVDData as (select  
    first_value(locale) over(partition by userid order by trans_dt asc ) as locale  
    , date_trunc('month', first_value(trans_dt)  
        over(partition by userid order by trans_dt asc))::date as cohort  
    , amt, userid, trans_dt  
from  
    cls.trans )  
  
select * from LTVDData where locale = 'U.S.' limit 100;
```

locale	cohort	amt	userid	trans_dt
U.S.	2016-05-01	23.98	1	2016-05-09
U.S.	2017-03-01	43.16	3	2017-03-05
U.S.	2017-03-01	43.16	3	2017-04-05
U.S.	2016-02-01	59.95	4	2016-02-28
U.S.	2016-01-01	99.95	6	2016-01-05
[...]				

- We can also have multiple tables defined:

```

with
    LTVDData as (select
        first_value(locale) over(partition by userid order by trans_dt asc ) as locale
        , date_trunc('month', first_value(trans_dt)
            over(partition by userid order by trans_dt asc))::date as cohort
        , amt, userid, trans_dt
    from
        cls.trans )
    , SubscribersFirst as (select distinct userid from
        (select userid, first_value( type ) over(partition by userid
            ORDER BY trans_dt asc, type asc ) as firsttype
        from cls.trans) as innerQ
    where firsttype = 'Sub')
select
    *
from
    SubscribersFirst
left join
    LTVDData
using(userid);

  userid  locale    cohort      amt  trans_dt
-----
      2  Canada    2018-08-01   12.99  2018-08-25
      3   U.S.     2017-03-01   43.16  2017-03-05
      3   U.S.     2017-03-01   43.16  2017-04-05
      5  Canada    2018-03-01   17.98  2018-03-09
      5  Canada    2018-03-01   17.98  2018-05-09
[...]
```

- The upside of using a CTE is that they can be much easier to read.
- There are two major downsides to using CTEs:
 1. Some databases do not support them (MySQL)
 2. In other databases they can act as optimization barriers. In particular, consider the following query:

```

with
  LTVData as (select
    first_value(locale)
      over(partition by userid order by trans_dt asc ) as locale
    , date_trunc('month', first_value(trans_dt)
      over(partition by userid order by trans_dt asc)) as cohort
    , amt, userid, trans_dt
  from
    cls.trans )
select
  *
from
  LTVData
where userid = 2;

```

locale	cohort	amt	userid	trans_dt
-----	-----	-----	-----	-----
Canada	2018-08-01 00:00:00+00:00	12.99	2	2018-08-25

In this example, it is clear that the filter `WHERE userid = 2` could be applied within the `LTVData` expression. However, it is not and the database will compute the entire `LTVData` before applying the filter, a costly choice. We will explore performance considerations in the next section.